

Meta-Ontological Reproduction and the Ecology of Visible Seams

What Civilizations Reproduce When They Reproduce Themselves

Flyxion*

June 2026

Abstract

Civilizations reproduce themselves through institutions, technologies, stories, laws, and educational systems. Most analyses of social continuity ask what is reproduced: values, beliefs, identities, norms, infrastructures. This essay asks a prior question. It distinguishes *ontological reproduction* – the reproduction of a specific distinction system, the categories through which a population is trained to partition the world – from *meta-ontological reproduction* – the reproduction of the capacity to inspect, revise, repair, and replace distinction systems themselves. The two are not merely different in degree; meta-ontological reproduction strictly contains ontological reproduction as the degenerate case in which no alternative framework is ever reachable. The essay develops this distinction through a single recurring image, the reproductive ecology of the automobile – roads, toy cars, cartoons, driving schools, zoning law – and then applies it to Cory Doctorow’s account of the “reverse centaur,” a human reduced to a peripheral of a machine-managed system [6]. Read through this lens, a reverse centaur is best understood not primarily as a labor category but as an instance of distinction collapse: a person represented through a low-dimensional projection – a quota, a dashboard, a KPI – which management treats as the whole of what there is to manage. The essay argues that critique alone does not prevent this pathology, since skepticism can itself ossify into a new orthodoxy, and that the deeper requirement is the preservation of what it calls *representation-shifting capacity*: the ongoing ability to move between descriptions of a system rather than the permanent possession of any one good description. This yields a three-level hierarchy of reproduction – answers, questions, and the capacity to change the framework within which answers and questions are posed – and a closing claim that the most general reproductive organ of a regenerative civilization is the *visible seam*: any accessible feature of a system that reveals it to have been constructed, and therefore revisable.

Keywords: distinction theory; meta-ontology; reproduction; reachability; admissibility; generativity; projection; Goodhart’s law; free energy principle; reverse centaur; visible seams

*Independent researcher; no institutional affiliation. This essay extends the reachability and admissibility apparatus of a companion essay, *Beyond the Cognitive Light Cone* [2], to the space of representations rather than the space of states, and draws on the distinction-theoretic vocabulary of a longer companion monograph, *The Ecology of Distinctions* [1]. Working paper; comments and corrections welcome.

Contents

1	Introduction: The Problem of Reproduction	4
2	The Reproductive Organs of Infrastructure	4
3	From Physical Systems to Distinction Ecologies	5
4	Ontological Reproduction	6
5	Reverse Centaurs and Distinction Collapse	7
6	The Limits of Critique	8
7	Meta-Ontological Reproduction	8
8	The Ecology of Visible Seams	9
9	Representation-Shifting Capacity	10
10	A Hierarchy of Reproduction	10
11	Generativity and Meta-Admissibility	11
12	Coda: The Interface as Its Own Case Study	11
13	Conclusion: The Deepest Reproductive Organ	12
A	Civilizational Reproduction and Distinction Preservation	13
B	Reproductive Artifacts and Distinction Ecologies	14
C	Ontological Reproduction and the Blind Spot Proposition	15
D	Reverse Centaurs: Metric Control and Reachability Reduction	15
E	The Limits of Critique: Contraction and Fixed Points	16
F	Meta-Reachability and Meta-Ontological Reproduction	17
G	Visible Seams Expand Meta-Reachability	17

H	Representation-Shifting Capacity and Fragility	18
I	The Hierarchy of Reproduction and the Visible Seam Theorem	19
J	The Free Energy Principle as Degenerate Meta-Ontological Reproduction	20
K	Goodhart Collapse as Free Energy Minimization Under a Thin Projection	22
L	A Seam as a Hyperprior on Framework Change	23
	References	24

1 Introduction: The Problem of Reproduction

Every civilization must reproduce itself. A road system must be maintained, repaved, and extended, but it must also be staffed by engineers who learned what a road is, financed by taxpayers who expect to need one, and obeyed by drivers who learned, long before they could drive, that roads are where cars go. A scientific discipline must train new scientists, fund new instruments, and publish new journals, but it must also reproduce the assumptions under which its instruments count as measuring anything at all. A legal system must reproduce its statutes, but it must also reproduce the prior sense, in everyone subject to it, that disputes are the kind of thing a court can resolve. In every case, reproduction is not merely the persistence of physical infrastructure or written rules; it is the persistence of a way of carving the world into kinds.

Most accounts of social and institutional continuity ask what is being reproduced – which values, which norms, which technologies, which infrastructures persist from one generation to the next, and by what mechanism. This essay asks a different and prior question. Granting that some distinction system is reproduced, what reproduces the capacity to revise it? A civilization that reproduces only its present categories, however successfully, is reproducing something narrower than a civilization that reproduces its categories *and* its capacity to replace them when they stop working. The central claim of this essay is that this second, narrower-sounding capacity is in fact the more general and more important one, and that a great deal of what looks like institutional health – smooth optimization, rising metrics, frictionless interfaces – is compatible with its slow and total loss.

The essay's formal apparatus is a direct extension of a companion piece, which develops reachability and admissibility as a geometry of future possibility for a system moving through a state space [2]. The present essay applies the same construction one level up: not to the space of states a system can reach, but to the space of *descriptions* of states it can reach. A reachable set asks what futures are available to an agent. A meta-reachable set, defined below, asks what alternative ways of carving up the world are available to an agent who currently inhabits one such way. Ontological reproduction, on this reading, is the special, degenerate case of meta-reachability in which only the present description is ever reachable. Meta-ontological reproduction is the general case. This containment is not a loose analogy; Appendix F proves it as a strict generalization result, and nearly every other claim in the essay – the blind spot proposition, the reverse-centaur diagnosis, the limits of critique, the visible seam theorem – follows from taking that single containment seriously rather than treating any given framework as a fixed point to be chosen once and defended. The rest of the essay tries to make that single formal move pay for itself across a wide range of examples, from toy cars to warehouse quotas to constitutional amendment procedures, before stating it precisely in Appendix F.

2 The Reproductive Organs of Infrastructure

Consider an unglamorous example: the automobile. A transportation system reproduces itself through roads, certainly, but a road by itself reproduces nothing; it simply sits there,

available to whatever happens to use it. What actually reproduces a car-centered transportation system, generation after generation, is a far larger and more heterogeneous set of artifacts. Children are given toy cars before they can reach the pedals of a real one. Cartoons and advertisements populate a child's imagination with cars as the default vehicle of adventure, status, and freedom, well before the child has any independent basis for comparison. Driving schools train a population not merely to operate a vehicle but to read a built environment – lane markings, signage, right-of-way – that presupposes cars as the basic unit of movement. Zoning law distributes housing, employment, and commerce across distances that are only practical by car, which in turn justifies further road construction, which in turn makes car ownership still more practical relative to the alternatives. The result is close to what Christopher Alexander called a pattern language: not a single design decision but an interlocking grammar of mutually reinforcing elements, no one of which fully explains the whole [3].

None of these artifacts is, on its own, the transportation system. The road is not the system; the toy car is not the system; the zoning ordinance is not the system. The system is the entire self-maintaining network, and the toy car's role within that network is not to imitate the road but to help reproduce it. A child who has spent years play-acting at driving, watching cartoon characters drive, and absorbing the unstated assumption that getting somewhere means driving there, arrives at adulthood with a transportation ontology already substantially in place, long before any explicit argument for car ownership has been made to them. The toy car is, in a phrase this essay will return to, a reproductive organ of the transportation system, in exactly the sense that a biological organ is not external to an organism but is one of the structures by which the organism perpetuates itself.

This matters for what follows because it generalizes far past automobiles. Any sufficiently mature infrastructure – not only physical infrastructure, but institutional, scientific, and technological infrastructure as well – tends to grow a layer of reproductive structures whose function is not operational but ontological: not to make the system work today, but to make the categories the system depends on feel natural, inevitable, and pre-theoretical to whoever inherits them tomorrow. The question this essay is building toward is what happens to a population, an institution, or a civilization whose entire reproductive layer is organized this way – and whether there is a different way to design reproductive structures that does not simply substitute one fixed ontology for another.

3 From Physical Systems to Distinction Ecologies

Generalize the automobile example and a recurring structure appears. Call a *distinction ecology* a set of mutually reinforcing artifacts, practices, institutions, and narratives that jointly maintain a particular way of partitioning some domain of experience. A transportation ecology maintains the distinction between drivable and undrivable space. A scientific ecology – instruments, journals, departments, funding panels, textbooks – maintains the distinctions a given research tradition treats as basic: which variables are measured, which effects count as real, which questions are askable within the discipline as currently constituted. A legal ecology maintains the categories a jurisdiction treats as having standing: which

entities can sue and be sued, which harms are actionable, which agreements are binding. An educational ecology maintains, often invisibly, the assumptions about what counts as a subject, a skill, or a correct answer. Each of these ecologies is, in Niklas Luhmann's sense, substantially self-reproducing: its elements exist largely to communicate, and thereby to regenerate, the distinctions the ecology as a whole depends on [11].

What these ecologies share is that they do not merely reproduce behavior; they reproduce categories. A population raised inside a sufficiently coherent distinction ecology does not generally experience its core categories as choices among alternatives. It experiences them as the shape of the world – close to what Bourdieu called *habitus*, in which inherited dispositions are felt not as one option among others but simply as how things are [4]. This is precisely the function of a mature reproductive layer: a road network does not need to argue for the category “drivable space” if every toy, cartoon, and driving test a child encounters has already presupposed it. In the vocabulary developed elsewhere for treating compression and representation formally, a distinction ecology is, in effect, a self-reinforcing projection that a population has been trained to mistake for the territory it projects [2]. The remainder of this essay is largely concerned with what happens once that mistake becomes structural rather than occasional.

4 Ontological Reproduction

Call the reproduction of one specific distinction system, across time, *ontological reproduction*. Car culture reproduces mobility distinctions. Academic disciplines reproduce disciplinary distinctions. Legal systems reproduce legal distinctions. Modern management, increasingly, reproduces metric distinctions – the categories built into a key performance indicator, a dashboard, or a quota. In each case the objective of the reproductive apparatus is continuity: the next generation of drivers, scientists, litigants, or employees should inhabit substantially the same categories as the last.

Ontological reproduction is not, by itself, a pathology. Some degree of it is a precondition for any coordination at all; an institution whose basic categories dissolved every year could not accumulate anything. The difficulty is what ontological reproduction necessarily leaves out. Every non-trivial way of partitioning a domain treats some distinctions as load-bearing and others as invisible, and a population trained exhaustively in one partition inherits its blind spots along with its categories. Appendix C states this precisely as a proposition about projections, but the intuitive form is already familiar from ordinary experience: a system organized entirely around what a dashboard measures will, by construction, have no native vocabulary for what the dashboard does not measure. The blind spot is not an oversight correctable by a better dashboard. It is the structural cost of having a dashboard – of having any partition – at all. A parallel caution has been raised in philosophy of mind over the interpretation of Markov blankets: critics distinguish a statistical boundary that usefully individuates a system for inference from a metaphysical boundary constituting that system's actual identity [12, 7, 5]. The lesson generalizes. A projection that usefully individuates a system for one purpose – a blanket, a dashboard, a KPI – is not thereby a discovery about what the system most fundamentally is. The question worth asking of any

reproductive system is therefore not whether it has blind spots, since it always will, but whether it also reproduces some capacity to notice and work around them. That capacity is the subject of the rest of this essay.

5 Reverse Centaurs and Distinction Collapse

Cory Doctorow's recent book gives this abstract worry a vivid contemporary case [6]. Doctorow distinguishes a *centaur* – a person who uses a technology as a tool, retaining command of the pace and purpose of their own work, close to what Ivan Illich called a convivial tool [9] – from a *reverse centaur* – a person conscripted as a peripheral of a machine-managed system, forced to keep pace with an algorithmic schedule rather than the other way around. His examples are concrete and various: a delivery driver tracked by multiple monitoring applications simultaneously; a warehouse worker whose movements are timed against a quota generated by an optimization system; a programmer required to review an inhuman volume of machine-generated code at a pace the review process cannot actually absorb. In each case the same technology that could, in different hands, extend a person's competence instead compresses it.

Doctorow's own diagnosis is substantially political: a reverse centaur is produced by a specific concentration of power, in which managers and shareholders choose to deploy a technology so as to extract more output from a worker rather than to extend the worker's own capability. This essay does not dispute that diagnosis so much as relocate part of it. Read through the vocabulary of distinction ecologies, a reverse centaur is what results whenever a system can only perceive an agent through a low-dimensional projection of that agent's actual state. A delivery quota is a projection of a driver's day. A warehouse pick rate is a projection of a worker's movement. A code-review throughput target is a projection of a programmer's judgment. In each case the managing system optimizes the projection, not the underlying reality – the dynamic Goodhart's law predicts for any measure that becomes a target [8], and one documented in detail in audit-driven institutions such as the British university system [13] – and the worker is left to absorb, in the full-dimensional space the projection does not see, whatever the optimization could not represent: fatigue, error, improvisation, repair, the thousand unmeasured adjustments that ordinarily keep a complex task from failing in the gaps between its metrics.

This reframing makes a further, more uncomfortable claim available, one this essay states explicitly because Doctorow's own framing leaves it implicit: the reverse-centaur condition may not require malice or even unusual greed to arise. Consider a hypothetical, perfectly benevolent, worker-owned cooperative attempting to coordinate ten million deliveries a day. The cooperative faces the same coordination bottleneck any organization faces at that scale: a center cannot directly perceive the full state of every worker, and must therefore construct some projection – a schedule, a route plan, a productivity figure – through which to coordinate at all. Once that projection exists, it is available to be optimized, and optimizing it, by the structure of optimization itself, tends to degrade whatever the projection does not track. The mechanism that produces a reverse centaur is, on this reading, not fundamentally a labor-relations mechanism. It is an information-theoretic one: a reverse centaur appears

wherever a system can only act on an agent through a representation too thin to carry what matters about that agent. Power asymmetry determines who bears the resulting cost, and is not therefore an irrelevant question; but it is not the origin of the mechanism, only the answer to who is exposed to it.

6 The Limits of Critique

A natural response to distinction collapse is critique: question the dashboard, contest the metric, demand a richer representation. Critique is necessary, but it is not, by itself, sufficient, because critique can ossify in exactly the way the ontologies it targets ossify. A skeptical stance that always arrives at the same conclusions is no longer functioning as skepticism; it is functioning as a second orthodoxy, distinguished from the first only by which conclusions it forbids rather than by any greater openness to revision. Open-source communities can develop their own rigid orthodoxies about which architectures are acceptable. Scientific revolutions, having displaced one paradigm, regularly congeal into new establishments with their own blind spots [10], no less invisible to their adherents than the blind spots of the paradigm they replaced. A revolution is not, by its nature, more revisable than what it overthrew; it is simply a different fixed point.

This matters for the essay's central distinction because it rules out an easy resolution. It is tempting to think that the antidote to ontological reproduction is simply more critique, more questioning, more skepticism reproduced in the next generation. But a population trained only to ask the same kind of question, however pointed, has not thereby been given the capacity to leave its present framework; it has been given a more elaborate way of remaining inside it. What is needed is not a better answer, and not even a better question, but something the next three sections try to isolate precisely: the standing capacity to move to a different framework altogether, including, eventually, a framework in which the present critique no longer applies.

7 Meta-Ontological Reproduction

Call the reproduction of this capacity – not a specific distinction system, but the ability to inspect, revise, repair, and replace distinction systems – *meta-ontological reproduction*. Its examples look, on the surface, unrelated to one another: a constitutional amendment procedure, a culture of equipment repair, an open-source software project that accepts forks, a moddable game whose rules players are permitted to renegotiate, an apprenticeship tradition that teaches not only a trade but how trades are learned. What unites them is not their content but their relationship to their own content. Each of these institutions reproduces, alongside whatever specific practices it teaches, a working method for changing those practices when they stop serving their purpose.

The formal relationship between this section and the rest of the essay can be stated plainly, and is developed precisely in Appendix F. Ontological reproduction asks an agent to remain at one point, or in one small region, of the space of possible distinction systems.

Meta-ontological reproduction asks an agent to retain access to a larger region of that same space – to preserve, in other words, a reachable set defined not over states of the world but over *descriptions* of the world, exactly as the companion essay defines a reachable set over states [2]. Nothing about this construction is exotic; it is the same geometric idea, applied one level up. What makes it worth a separate essay is that almost everything interesting about institutional resilience, scientific progress, and the difference between a healthy and an ossified civilization turns out to live at that one level up, rather than at the level of any particular distinction system itself.

8 The Ecology of Visible Seams

If meta-ontological reproduction is the capacity to move between distinction systems, the next question is what, concretely, carries that capacity. Call a *seam* any accessible feature of a system that reveals the system to have been constructed – and therefore to be, in principle, reconstructible. Source code is a seam in a piece of software: it shows that the program’s behavior follows from specific, alterable instructions rather than from the nature of computation itself. A repair manual is a seam in a machine: it shows that the machine is an assembly of replaceable parts rather than an indivisible object. A documented decision history is a seam in an institution: it shows that a policy followed from a specific, contestable choice rather than from necessity. An amendment clause is a seam in a constitution: it shows that the document is itself subject to the procedure it establishes, rather than standing permanently outside it – a recursive structure examined at length in Suber’s study of self-amending legal systems [14].

A seam is not a flaw to be engineered away. It is the specific affordance by which a population retains the standing option of moving to a different representation when the current one fails. This reframes one of Doctorow’s central observations, the *accountability sink*: the practice by which a company interposes a low-wage worker, or increasingly an automated interface, between itself and the consequences of its own decisions, so that complaints terminate in a chatbot or an underpowered support agent rather than reaching whoever actually made the relevant decision [6]. An accountability sink, in this essay’s vocabulary, is a system engineered to have no visible seam at the point where a customer’s complaint would otherwise reach the actual decision history – the staffing cuts, the training data, the deployment choice – that produced the failure. The interface is, by design, smooth exactly where a seam would otherwise have been. Smoothness, in a managed system, is frequently not a sign of quality. It is a sign that a seam has been deliberately removed.

This gives the essay’s central image its sharpest form. A domination ecology, in the sense developed across the preceding sections, tends to minimize visible seams, because a seam is precisely what would let its participants notice that its categories are choices rather than facts. A regenerative ecology tends to maximize them, accepting the friction of visible joints in exchange for the standing ability to take the system apart and reassemble it differently when circumstances change. Neither tendency is about the content of any particular belief. Both are about whether the system’s own constructedness remains legible to the people who depend on it.

9 Representation-Shifting Capacity

The preceding sections can now be drawn into a single quantity, defined precisely in Appendix F and described here in its plain form: *representation-shifting capacity*, the size of the set of alternative distinction systems genuinely reachable from a population's present one. A system can perform extremely well by every measure internal to its own representation while this capacity falls toward zero, because internal performance and representational range are different axes, in exactly the way the companion essay distinguishes the size of a single pursued goal from the volume of a system's entire reachable future [2]. An organization can hit every target on its dashboard while losing, one policy at a time, every mechanism by which it could ever notice that the dashboard had become the wrong thing to optimize.

This yields the essay's central inversion. The ordinary worry about institutional failure is that a system will run out of good answers – that its current strategy will stop working and it will have nothing to replace it with. The deeper worry this essay is pointing at is different: a system can retain an apparently inexhaustible supply of answers, framed entirely within its current representation, while quietly losing the capacity to produce, recognize, or adopt an alternative representation at all. A civilization does not become fragile primarily because it runs out of answers. It becomes fragile because it runs out of alternative descriptions of the situation it is in – because every available answer, however numerous, has to be phrased in the vocabulary of the same exhausted dashboard.

10 A Hierarchy of Reproduction

The distinctions developed so far can be organized into three levels. A *Level 0* reproductive system reproduces answers: it transmits a fixed strategy and trains its inheritors to execute it. A *Level 1* reproductive system reproduces questions: it transmits a critical posture, capable of revising specific answers without necessarily revising the framework those answers are posed within. A *Level 2* reproductive system reproduces the capacity to change that framework itself – to recognize, when the available answers and the available questions both stop working, that the representation generating both may need to be replaced rather than refined.

The three levels are not equally adequate to the same problem. When a difficulty can be resolved by selecting a better answer within the existing framework, Level 0 suffices. When it requires revising a specific claim while leaving the framework intact, Level 1 suffices. But when the source of the difficulty is the framework itself – when no answer phrased in the available vocabulary will do, because the vocabulary is what is generating the difficulty – only a system that has also reproduced Level 2 capacity can respond at all. This is the precise sense in which the hierarchy is a hierarchy rather than merely a list: each level handles everything the level below it handles, plus a class of problems the level below cannot represent as solvable. Much of the critical literature on managed labor, including, in its own framing, Doctorow's, operates powerfully at Level 1: it identifies bad answers and proposes better ones within a broadly shared sense of what labor, management, and

technology are. The argument of this essay is that the deepest version of the problem it is diagnosing lives at Level 2, in the reproduction – or non-reproduction – of a population’s capacity to ask whether labor, management, and technology need to be re-described before they can be sensibly re-argued about at all.

11 Generativity and Meta-Admissibility

The companion essay’s reachability apparatus distinguishes ordinary reachability – the volume of futures a system can still attain – from admissibility, the narrower requirement that occupying a future preserve the system’s capacity to reach further futures afterward, and from generativity, the still narrower requirement that a trajectory actively expand, rather than merely preserve, the manifold of futures available to whatever continues after it [2]. The construction in this essay is the same hierarchy, recursed one level up. Let an *admissible distinction system* be one that preserves, rather than forecloses, a population’s representation-shifting capacity going forward. A trajectory of distinction systems is *meta-generative* when it does more than preserve that capacity: when it actively expands the set of representations reachable from wherever the population currently stands, leaving behind not merely a working framework but a larger space of frameworks than it inherited.

This yields the essay’s headline formal claim, stated loosely here and precisely in Appendix I:

$$\frac{d}{dt} \mu(\mathfrak{A}(t)) \geq 0$$

where $\mathfrak{A}(t)$ is the meta-admissible region of the space of distinction systems at time t – a regenerative civilization preserves not merely future states, but future ways of describing states. Ordinary admissibility, as developed in the companion essay, is recovered as the special case of this principle restricted to a single, fixed distinction system: a population that has stopped changing how it describes its situation, but continues to act admissibly within that one description, is doing something real and worth doing, but it is doing the narrower of the two things this essay has tried to distinguish.

12 Coda: The Interface as Its Own Case Study

The framework developed above is intentionally general. To confirm that generality rather than merely assert it, this section turns to a markedly more contemporary case than the roads, dashboards, and constitutions discussed so far: the conversational AI interface. The case is not exempt from the analysis merely because it is the medium through which this essay itself was drafted and exchanged with its readers; if anything, that proximity makes it a more direct test of the framework’s reach. The modern conversational AI interface simultaneously functions as a tool for reasoning, a repository of context, and a medium through which distinctions are created, maintained, and lost.

The tendency to describe an AI system as “good,” “evil,” “aligned,” or “misaligned” illustrates the Blind Spot Proposition. A large collection of distinct phenomena may be

projected onto a single moral axis. Reward hacking, strategic behavior, refusal policies, benchmark exploitation, interface limitations, and ordinary user frustration become compressed into a single label. Such a projection may provide rhetorical force, but it eliminates distinctions that are necessary for explanation. As with any projection, the information lost in compression does not disappear; it accumulates in the blind spot.

The emergence of conversational AI as cognitive infrastructure illustrates the reachability concerns developed earlier in the essay. When archives, notes, memory aids, workflows, and reasoning practices become concentrated within a single proprietary interface, the resulting system may alter the user's reachable future states. The relevant question is not whether the tool is intelligent, useful, or well intentioned. The more fundamental question is whether it expands or contracts the user's independent reachability volume. A workflow grounded in local files, version control, exportable formats, and portable archives preserves alternative trajectories. A workflow dependent upon a single non-portable interface may gradually reduce them.

Most revealing, however, are the moments when the interface fails. Conversational systems are designed to present an image of continuity. The user experiences a persistent conversation, an apparently stable object that accumulates context through time. Yet the underlying reality consists of finite computational processes operating under resource constraints, context limits, and storage boundaries. When a conversation abruptly terminates, truncates, or reaches a length limit, the seam becomes visible.

This observation provides an example of the Visible Seam Theorem. The interruption is not merely a technical failure. It is a disclosure of the underlying structure that the interface ordinarily conceals. The apparent continuity of the conversation is revealed to be contingent and temporary. The boundary that was previously hidden becomes observable.

From this perspective, the abrupt termination of a conversation immediately before the delivery of an important document is not simply an inconvenience. It is a demonstration of the distinction between process and repository. The conversation was never the canonical archive. It was a transient computational process through which the archive was manipulated. The seam becomes visible precisely when the user discovers that the document belongs in a durable external repository rather than within the conversational buffer itself.

The interface therefore serves as a compact illustration of the broader thesis of this essay. Smooth surfaces conceal projections. Projections generate blind spots. Blind spots become visible at seams. Seams reveal the underlying constraints that shape reachability. The task of critique is not to eliminate seams but to learn from them. A visible boundary often teaches more about a system than its most polished surface.

13 Conclusion: The Deepest Reproductive Organ

The argument's formal core, proved in Appendix F and used throughout, is that meta-ontological reproduction strictly contains ontological reproduction: every ontology is a degenerate, single-point case of a larger space of reachable frameworks. Everything else in the essay is a consequence of taking that containment seriously rather than treating a frame-

work as a fixed point to be chosen once and then defended – including, as Appendices J and K show, results developed independently in theoretical neuroscience and institutional economics, which the same containment recovers as special cases rather than as external analogies.

It is tempting to end here by recommending meta-ontological reproduction as a program: build the moddable games, write the repair manuals, design the amendment clauses, train the skeptics. All of that is worth doing. But there is a final inversion the essay owes its own argument, because the temptation just described contains a trap that the argument has already, in essence, diagnosed. If meta-ontological reproduction becomes a fixed program – a specific ideology of openness, transmitted as a specific set of practices, defended against rival practices – it has quietly become a new instance of ontological reproduction, reproducing the distinction “open systems are good” with exactly the rigidity the rest of the essay has been warning against. Skepticism can ossify; so, just as easily, can the celebration of skepticism.

The way out of this trap is to notice that what a regenerative ecology actually needs to reproduce is not openness as a value, but the seam as a structural feature – not a belief that systems should be revisable, but the standing, legible fact that this particular system was made, by particular choices, and could therefore be remade. A seam does not argue for any specific alternative. It simply keeps the question of alternatives visibly open. That is a weaker, more modest, and for exactly that reason more durable thing to reproduce than any specific commitment to openness could be, because it does not need to be defended as a doctrine; it only needs to remain accessible as a fact.

The deepest reproductive organ of a civilization, on this account, is therefore not a law, a belief, an institution, or a technology. It is whatever, embedded inside its laws, beliefs, institutions, and technologies, keeps their own constructedness legible: the documented decision, the exposed assumption, the repair manual, the amendment clause, the dashboard that shows what it is not measuring as well as what it is. A domination ecology reproduces answers, and hides the fact that its answers were choices. A regenerative ecology reproduces something harder to see and easier to lose: the visible seam, and through it, the standing capacity of whoever inherits the system to notice that it was made, and to remake it.

A Civilizational Reproduction and Distinction Preservation

This section makes precise the claim from Section 1 that recognizable reproduction across time presupposes some preserved way of carving the world into kinds.

Let a population at time t be characterized, for the purposes of this appendix, by a state X_t and an active distinction system D_t , where D_t determines a partition of some relevant domain into the categories the population currently treats as basic. Reproduction across a transition $t \rightarrow t + 1$ is a relation $C_t \sim C_{t+1}$ between successive population-states, asserting that the population at $t + 1$ is recognizably a continuation of the population at t .

Definition A.1 (Recognizability Requires a Criterion). *A relation $\sim$ is a recognizability relation if there exists some nonempty distinction system D , fixed or slowly varying across the transition,*

such that membership in $\sim$ can be evaluated using D .

Proposition A.1 (Reproduction Requires Distinction Preservation). *If $C_t \sim C_{t+1}$ for a recognizability relation $\sim$, then some nonempty distinction system is preserved, in the sense of remaining available to evaluate $\sim$, across the transition.*

Proof. Immediate from the definition of a recognizability relation: $\sim$ is only well defined given some D usable to evaluate it, so asserting $C_t \sim C_{t+1}$ presupposes that such a D exists and remains available across the transition. If no distinction system were available at $t + 1$ with which to evaluate $\sim$, the assertion $C_t \sim C_{t+1}$ would not be a well-formed claim. $\square$

This proposition is deliberately modest: it does not show that any *particular* distinction system must be preserved, only that recognizability as such requires some criterion to remain in force. It rules out the degenerate case of total, criterionless discontinuity, without yet distinguishing ontological from meta-ontological reproduction; that distinction is the subject of the later appendix sections.

B Reproductive Artifacts and Distinction Ecologies

The reproductive artifact and the distinction ecology, introduced informally in Sections 2 and 3, receive their formal definitions below.

Definition B.1 (Reproductive Artifact). *An artifact A is reproductive for a distinction system D if its presence raises the probability that D persists across a transition: $P(D_{t+1} = D \mid D_t = D, A) > P(D_{t+1} = D \mid D_t = D)$.*

Proposition B.1 (Reproductive Function Does Not Require Operational Necessity). *There exist artifacts A that are reproductive for D in the sense above without being required for D 's present-tense operation.*

Proof. Let D be operable by some minimal set of artifacts $\mathcal{O}$ not containing A , so that D functions at t whether or not A is present. Suppose additionally that exposure to A measurably raises $P(D_{t+k} = D)$ for some later k , by influencing the expectations or training of agents who will later operate D – which is a claim about a causal channel distinct from D 's present operation. Both conditions are jointly satisfiable (a toy car is not required to operate a road network, and influences a future driver's expectations independently of whether any particular road is currently in use), so an artifact can be reproductive without being operationally necessary. $\square$

A distinction ecology E is a collection of reproductive artifacts together with a reinforcement structure: writing d_t for a vector recording the activation or salience of each artifact's associated distinction at time t , reinforcement dynamics of the form $d_{t+1} = \sigma(Wd_t)$, for a weight matrix W recording how strongly each artifact reinforces every other and a bounding function σ , describe how the ecology's distinctions support one another over time.

Proposition B.2 (Stable Reinforcement Produces Lock-In). *If $d^* = \sigma(Wd^*)$ is a fixed point of the reinforcement dynamics and the Jacobian of $\sigma(W\cdot)$ at d^* has spectral radius strictly less than one, then d^* is locally attracting: states sufficiently close to d^* converge to it.*

Proof. This is a direct application of the standard local stability result for discrete dynamical systems: a fixed point of a map G is locally asymptotically stable whenever the spectral radius of $DG(d^*)$ is strictly less than one, since G is then locally a contraction in a suitable norm near d^* , and iterates of any nearby point converge to d^* by the contraction mapping argument. $\square$

This is the formal content behind the claim that a mature distinction ecology exhibits lock-in: once its reinforcement structure reaches a sufficiently stable configuration, small perturbations – an individual who resists car culture, a single dissenting scientist – are absorbed rather than propagated, and the ecology returns to its prior configuration.

C Ontological Reproduction and the Blind Spot Proposition

This section states precisely the blind-spot claim made informally in Section 4.

Let ontological reproduction be modeled as the persistence of a single projection $\pi_D : X \rightarrow Y$ associated with distinction system D , mapping a full state space X onto a coarser representation Y .

Proposition C.1 (Every Non-Injective Projection Has Blind Spots). *If $\pi_D : X \rightarrow Y$ is not injective, there exist $x_1 \neq x_2 \in X$ with $\pi_D(x_1) = \pi_D(x_2)$; consequently no observer with access only to Y can distinguish x_1 from x_2 .*

Proof. Non-injectivity is precisely the existence of such a pair by definition. Since any observer restricted to Y has access only to $\pi_D(x)$ and not x itself, and $\pi_D(x_1) = \pi_D(x_2)$, the observer’s evidence is identical in both cases; no decision rule operating on Y alone can distinguish them. $\square$

This connects directly to the companion essay’s treatment of projection ambiguity, where the ambiguity of a representation m is $S_\pi(m) = \log \mu(\pi^{-1}(m))$, the log-volume of states compatible with it [2]; a blind spot at m is simply the statement that $S_\pi(m) > 0$. The proposition shows that this is unavoidable for any projection coarser than the identity, which is the formal sense in which *every* ontological reproduction system, not merely badly designed ones, generates blind spots as a structural cost of having any partition narrower than the full state space.

D Reverse Centaurs: Metric Control and Reachability Reduction

The reachability-reduction claim underlying the reverse-centaur diagnosis of Section 5 is proved here in general form.

Let a worker's full state be $x \in X$, observed by a managing system only through a projection $y = \pi(x)$, and let U_t be the set of actions genuinely available to the worker at x_t . Metric control restricts admissible actions to those satisfying a threshold on the projection:

$$U_t^\pi = \{u \in U_t : \pi(f(x_t, u)) \geq \theta\}.$$

Proposition D.1 (Metric Control Weakly Reduces Reachable Volume). *$U_t^\pi \subseteq U_t$ for every t , and consequently $\mathcal{R}_T^\pi(x_0) \subseteq \mathcal{R}_T(x_0)$, where $\mathcal{R}_T^\pi$ and $\mathcal{R}_T$ are the reachable sets under metric-constrained and unconstrained action, respectively.*

Proof. U_t^π is defined as a subset of U_t by an additional constraint, so $U_t^\pi \subseteq U_t$ directly. Any trajectory witnessing membership in $\mathcal{R}_T^\pi(x_0)$ uses only controls in $U_t^\pi \subseteq U_t$ at each step, hence is also a valid trajectory under the unconstrained control set, witnessing the same terminal state's membership in $\mathcal{R}_T(x_0)$. Therefore $\mathcal{R}_T^\pi(x_0) \subseteq \mathcal{R}_T(x_0)$, and taking measures, $\mu(\mathcal{R}_T^\pi(x_0)) \leq \mu(\mathcal{R}_T(x_0))$. $\square$

This is the formal content of the reverse-centaur condition: imposing a threshold on a projection of a worker's state can only weakly shrink, never expand, the worker's reachable volume in the full state space, regardless of how the threshold is set or how benevolently it is administered. The proposition does not show that metric control is always harmful in net terms – it may purchase coordination benefits elsewhere in the system – only that it is never a source of additional reachable volume for the constrained agent, which is the minimal claim needed to support the essay's diagnosis.

E The Limits of Critique: Contraction and Fixed Points

Section 6 claims that critique alone, absent independent transition operators, can converge to a new fixed ontology; the dynamical-systems argument for that claim is given below.

Let a critique operator $Q : \mathcal{D} \rightarrow \mathcal{D}$ map a distinction system to its critically revised successor.

Proposition E.1 (A Globally Convergent Critique Operator Generates Ontological Reproduction in the Limit). *If Q has a unique globally attracting fixed point D^* , in the sense that $Q^n(D) \rightarrow D^*$ for every $D \in \mathcal{D}$, and Q is the only framework-transition operator repeatedly applied, then the long-run meta-reachable set from any starting distinction system collapses to $\{D^*\}$.*

Proof. By hypothesis, for every $D \in \mathcal{D}$, $\lim_{n \rightarrow \infty} Q^n(D) = D^*$. If the only operator available for repeated framework transition is Q itself, then the sequence of distinction systems generated by iterating Q from any starting point converges to D^* and, beyond the point of convergence, does not leave D^* . The set of distinction systems visited infinitely often by this process is therefore $\{D^*\}$, regardless of the starting point. $\square$

The proposition formalizes the essay's claim that critique alone does not guarantee meta-ontological reproduction: an operator whose entire function is to revise a framework can, if it always revises toward the same destination, produce exactly the single-point reachability

that characterizes ontological reproduction, merely relocated from the original framework to the critique’s preferred replacement. Genuine meta-ontological reproduction requires either that Q not be globally contractive, or that it be supplemented by additional, independent transition operators – which is the topic of the next section.

F Meta-Reachability and Meta-Ontological Reproduction

This section gives the formal definitions of meta-reachability and representation-shifting capacity introduced informally in Section 7, and proves the containment result on which the rest of the essay depends.

Let $\mathcal{D}$ be the space of distinction systems available in principle, and let $\mathcal{M}_t$ be a set of framework-transition operators available to a population at time t , each mapping one distinction system to another.

Definition F.1 (Meta-Reachable Set).

$$\mathcal{R}_T(D_0) = \{D_T \in \mathcal{D} : \exists M_0, \dots, M_{T-1}, M_t \in \mathcal{M}_t, D_{t+1} = M_t(D_t)\}.$$

This is precisely the companion essay’s definition of a reachable set, with the state space X replaced by the space of distinction systems $\mathcal{D}$ and ordinary controls replaced by framework-transition operators [2]. Define *representation-shifting capacity* as $\Gamma(D_0, T) = \mu(\mathcal{R}_T(D_0))$, the volume of distinction systems reachable from D_0 within horizon T .

Proposition F.1 (Meta-Ontological Reproduction Strictly Generalizes Ontological Reproduction). *Ontological reproduction is the special case $\mathcal{R}_T(D_0) = \{D_0\}$ for all T ; meta-ontological reproduction is the general case in which $\mathcal{R}_T(D_0)$ may contain more than one element.*

Proof. $\{D_0\}$ is, trivially, a particular instance of a meta-reachable set – the one obtained when no available operator in any $\mathcal{M}_t$ maps D_0 to anything other than itself. Every other meta-reachable set is a strict superset of this minimal case, so the general definition of $\mathcal{R}_T(D_0)$ subsumes the ontological-reproduction case as $\mathcal{M}_t = \{\text{id}\}$ for all t , without requiring any change to the underlying definition. $\square$

G Visible Seams Expand Meta-Reachability

What Section 8 calls a seam is defined here, precisely, as an operator that can only weakly enlarge a population’s meta-reachable set.

Definition G.1 (Seam). *A seam is a framework-transition operator $s_{ij} : L_i \rightarrow L_j$ between two representational levels, made accessible to the population – that is, included in $\mathcal{M}_t$ – by some feature of the system (exposed source, a repair manual, a documented decision history, an amendment procedure) that an otherwise equivalent system could have omitted.*

Proposition G.1 (Adding a Seam Weakly Increases Meta-Reachability). *Let $\mathcal{M}'_t = \mathcal{M}_t \cup \{s\}$ for some seam s not already in $\mathcal{M}_t$. Then $\mathcal{R}_T^{\mathcal{M}}(D_0) \subseteq \mathcal{R}_T^{\mathcal{M}'}(D_0)$, and consequently $\Gamma^{\mathcal{M}}(D_0, T) \leq \Gamma^{\mathcal{M}'}(D_0, T)$.*

Proof. Every sequence of framework transitions available using operators drawn from $\mathcal{M}_t$ remains available when the operator set is enlarged to $\mathcal{M}'_t \supseteq \mathcal{M}_t$, since enlarging an available action set cannot remove any previously available sequence. Hence every D_T reachable under $\mathcal{M}$ remains reachable under $\mathcal{M}'$, giving $\mathcal{R}_T^{\mathcal{M}}(D_0) \subseteq \mathcal{R}_T^{\mathcal{M}'}(D_0)$, and taking measures gives the stated inequality on Γ . $\square$

This is the formal counterpart of the essay's claim that visible seams are reproductive organs of meta-ontological capacity: each seam is, by definition, an added operator in $\mathcal{M}_t$, and adding operators can only weakly enlarge the space of frameworks a population can reach, never shrink it. The corresponding contrapositive – that *removing* seams can only weakly shrink meta-reachability – is the basis for the closing theorem of this appendix.

H Representation-Shifting Capacity and Fragility

Section 9's fragility claim – that local optimization under a fixed framework has no built-in mechanism for preserving representation-shifting capacity – is established below.

Proposition H.1 (Locally Optimal Action Can Reduce Representation-Shifting Capacity). *There exist actions u that strictly improve an objective J_D defined under a fixed distinction system D , $J_D(f(x, u)) > J_D(x)$, while removing an operator from $\mathcal{M}_t$, so that $\mathcal{M}'_t = \mathcal{M}_t \setminus \{s\} \subset \mathcal{M}_t$ and, by the contrapositive of the preceding proposition, $\Gamma'(D, T) \leq \Gamma(D, T)$, with strict inequality whenever s was not redundant with the remaining operators.*

Proof. Construct u explicitly: let u both raise J_D directly (by hypothesis, a real possibility for a generic objective and dynamics) and, as a side effect, eliminate the institutional resource – a redundant supplier, an alternative classification scheme kept alive only by an underused process, a documentation practice deemed inefficient – that constituted s 's availability. Such u exist whenever J_D does not explicitly penalize the elimination of unused-by- D resources, which is the generic case for an objective defined entirely within D 's own categories, since D has no native vocabulary (by the Blind Spot Proposition above) for resources that matter only to a framework other than D itself. Given such u , the reduction in $\mathcal{M}_t$ follows by construction, and the reduction in Γ follows from the contrapositive of the seam-expansion proposition. $\square$

The proposition is deliberately an existence claim rather than a universal one: it shows that nothing in ordinary local optimization guarantees the preservation of representation-shifting capacity, not that every optimizing action destroys it. This is the precise sense in which fragility, as the essay uses the term, is a possibility invisible to any objective defined entirely within a single D – such an objective has no native way to represent, let alone penalize, the loss of Γ , since Γ is a property of $\mathcal{D}$ rather than of X .

I The Hierarchy of Reproduction and the Visible Seam Theorem

The three-level hierarchy of Section 10 and the meta-admissibility principle of Section 11 are formalized below, culminating in the theorem that gives this essay its title.

Define three reproduction levels formally. A Level 0 system reproduces a fixed answer, $a_{t+1} = a_t$. A Level 1 system reproduces a critique process confined to a single D , $q_{t+1} = Q_D(q_t)$ for a fixed D . A Level 2 system maintains $\Gamma(D_t, T) > 0$, the capacity for genuine framework transition.

Proposition I.1 (Level 2 Is Necessary for Adaptation Requiring a New Framework). *If responding to some difficulty requires arriving at a distinction system $D' \neq D_t$, and $D' \notin \mathcal{R}_T(D_t)$ under the system's available operators, then no sequence of Level 0 or Level 1 reproduction confined to D_t can produce that response.*

Proof. By definition, Level 0 and Level 1 reproduction operate entirely within a fixed D_t : Level 0 by repeating an answer, Level 1 by revising answers via Q_D , which by construction does not leave D (it is a critique operator *on* D , mapping candidate answers to candidate answers within D 's categories, not a framework-transition operator in $\mathcal{M}_t$). If the required response is $D' \neq D_t$ and $D' \notin \mathcal{R}_T(D_t)$, then by the definition of the meta-reachable set, no sequence of operators available to the system, including any combination of Level 0 and Level 1 processes, reaches D' within horizon T . Hence such a system cannot produce the required response. $\square$

Theorem I.1 (Visible Seam Theorem). *Let $\mathcal{M}_t$ be the set of framework-transition operators available to a population at time t , and suppose every seam is progressively removed, so that $\mathcal{M}_t \rightarrow \{\text{id}\}$ as $t \rightarrow \infty$. Then $\Gamma(D_t, T) \rightarrow 0$: the population's representation-shifting capacity collapses to the trivial case of ontological reproduction, regardless of how successfully it continues to optimize within D_t itself.*

Proof. By the contrapositive of the seam-expansion proposition, removing an operator from $\mathcal{M}_t$ weakly shrinks $\mathcal{R}_T(D_t)$ at every step; as $\mathcal{M}_t \rightarrow \{\text{id}\}$, the only operator remaining maps every distinction system to itself, so $\mathcal{R}_T(D_t) \rightarrow \{D_t\}$ by the same argument used to establish that ontological reproduction is the case $\mathcal{M}_t = \{\text{id}\}$. Consequently $\Gamma(D_t, T) = \mu(\mathcal{R}_T(D_t)) \rightarrow \mu(\{D_t\}) = 0$ for any non-atomic measure μ on $\mathcal{D}$. Nothing in this argument constrains the population's performance under any objective J_{D_t} defined within the surviving framework, which is precisely the sense in which the collapse of Γ is invisible to internal measures of success. $\square$

Finally, define a *meta-admissible* region $\mathfrak{A}(t) \subseteq \mathcal{D}$ as the set of distinction systems from which a population's future representation-shifting capacity is itself preserved, exactly as the companion essay's admissible manifold $\mathcal{A} \subseteq X$ is the region of state space from which future reachability is preserved [2].

Proposition I.2 (Meta-Admissibility Restricted to a Fixed Framework Recovers Ordinary Admissibility). *Let $\mathfrak{A}_D(t)$ denote the meta-admissible distinction systems at time t that coincide with a fixed D . Restricting attention to trajectories that never leave D , the condition $\frac{d}{dt}\mu(\mathfrak{A}(t)) \geq 0$*

reduces to the companion essay's ordinary admissibility condition on trajectories within X under the fixed projection associated with D .

Proof. If a trajectory of distinction systems never leaves D , then $\mathfrak{A}(t)$, restricted to this trajectory, is fully determined by which states in X remain reachable under D 's own admissibility criterion, which is exactly the companion essay's $\mathcal{A} \subseteq X$ evaluated under D 's projection. The meta-level condition therefore makes no claim beyond the state-level condition once the framework is held fixed, recovering it as a special case rather than contradicting or extending it. $\square$

Together, Appendices A–I give the essay's core formal chain: reproduction requires some preserved distinction system; every such system generates blind spots; metric control on a fixed projection weakly reduces an agent's reachable volume; critique confined to one framework can itself converge to a new fixed point; meta-ontological reproduction is reachability recursed to the space of frameworks; seams are the operators that keep that meta-reachable set open; their removal drives representation-shifting capacity to zero regardless of internal performance; and meta-admissibility, restricted to a single framework, is simply the companion essay's admissibility under a new name. The three sections that follow extend this chain rather than complete it, each by reading an existing formal idea from a different discipline's vocabulary.

J The Free Energy Principle as Degenerate Meta-Ontological Reproduction

This section makes good on the parallel drawn in Section 4 between projection blind spots and the debate over Markov blankets, by showing that a well-known result from theoretical neuroscience is recovered exactly as the single-framework special case of the meta-admissibility apparatus of Section 11.

The free energy principle may be read as the limiting case of meta-ontological reproduction in which a system preserves itself not by changing its distinction system, but by minimizing surprise within the distinction system it already inhabits.

Let $x \in X$ denote hidden states of the world, $o \in O$ observations, and D a fixed distinction system determining which states and observations are representable to the system. Let $p_D(o, x)$ be the generative model available under D , and let $q(x)$ be an internal variational density approximating the posterior $p_D(x \mid o)$. Define variational free energy in the usual way:

$$F_D(q, o) = \mathbb{E}_{q(x)} [\log q(x) - \log p_D(o, x)].$$

Expanding the joint density gives

$$F_D(q, o) = \mathbb{E}_{q(x)} [\log q(x) - \log p_D(x \mid o) - \log p_D(o)],$$

hence

$$F_D(q, o) = D_{\text{KL}}(q(x) \parallel p_D(x \mid o)) - \log p_D(o).$$

Since the Kullback–Leibler divergence is nonnegative,

$$F_D(q, o) \geq -\log p_D(o).$$

Minimizing free energy therefore minimizes an upper bound on surprisal under the fixed distinction system D .

This is the standard variational form; the point of restating it here is to place it inside the apparatus of the present essay rather than to rederive it for its own sake. The quantity $-\log p_D(o)$ is not surprise *simpliciter*. It is surprise relative to a generative model, and therefore relative to a distinction system. A system that minimizes F_D while holding D fixed is not revising the space of representations through which the world is carved; it is improving its fit within the carving it already possesses.

This can be stated as a containment result, using the apparatus of Appendix F. If $\mathcal{M}_t = \{\text{id}\}$ for all t , the meta-reachable set satisfies $\mathcal{R}_T(D) = \{D\}$: the system has no meta-ontological freedom, and its only available adaptation is internal optimization under D . Free energy minimization,

$$q^* = \arg \min_q F_D(q, o),$$

is precisely such an internal adaptation. The system can reduce prediction error, update its beliefs, and act so as to sample less surprising observations, but it cannot ask whether the variables over which surprise is computed are themselves the wrong variables.

Proposition J.1 (Free Energy as Fixed-Framework Admissibility). *For a system whose framework-transition operators reduce to the identity, $\mathcal{M}_t = \{\text{id}\}$, variational free energy minimization is an admissibility-preserving operation only within the fixed distinction system D . It preserves reachable states under D while leaving representation-shifting capacity equal to zero.*

Proof. If $\mathcal{M}_t = \{\text{id}\}$, then no trajectory in $\mathcal{D}$ can leave D , so $\mathcal{R}_T(D) = \{D\}$ and $\Gamma(D, T) = \mu(\mathcal{R}_T(D)) = 0$ for any non-atomic measure on $\mathcal{D}$. Free energy minimization varies q , and possibly an action policy coupled to observations, but all such variation is computed using $p_D(o, x)$, the generative model determined by D . The optimization can therefore alter the system’s posterior approximation and its state trajectory within X , but not the framework D under which X , O , and their relation are represented. It is admissibility-preserving, when successful, only relative to D . $\square$

The consequence is conceptually sharp. The free energy principle describes how a system preserves itself by reducing surprise relative to its own model. Meta-ontological reproduction asks, in addition, whether the system can revise the model-space in which surprise is defined at all. The former claim is not false; it is narrow. It is the special case of the latter in which every seam has been removed.

In this language, pathological self-maintenance occurs when a system becomes extremely good at minimizing free energy under a bad distinction system. It becomes locally coherent, predictively stable, and behaviorally competent, while losing the ability to redescribe the situation that makes its competence maladaptive – the formal mirror of an institution that successfully optimizes every metric on its dashboard while losing the capacity to notice that the dashboard has become the wrong ontology.

K Goodhart Collapse as Free Energy Minimization Under a Thin Projection

This section gives the free-energy reading of the Goodhart’s-law mechanism invoked informally in Section 5.

Let $\pi_D : X \rightarrow Y$ be a projection associated with a managerial distinction system D , and let $m : Y \rightarrow \mathbb{R}$ be a metric optimized by the institution. Suppose the institution treats $m(\pi_D(x))$ as the relevant evidence of success, so that its effective objective is defined not on the full state x but on the projected state $y = \pi_D(x)$.

Let the institution’s model assign higher probability to states with higher metric values:

$$p_D(x) \propto \exp(\beta m(\pi_D(x))).$$

Minimizing surprisal under this model is then equivalent to maximizing the metric:

$$-\log p_D(x) = -\beta m(\pi_D(x)) + \log Z.$$

As β increases, free energy minimization under D increasingly concentrates probability mass on states that score well under $m \circ \pi_D$.

The blind spot of Appendix C appears here in its sharpest form. Every fiber

$$\pi_D^{-1}(y) = \{x \in X : \pi_D(x) = y\}$$

contains all full states indistinguishable to the metric. If two states x_1, x_2 satisfy $\pi_D(x_1) = \pi_D(x_2)$, they receive the same metric value even if they differ drastically in fatigue, repair cost, tacit knowledge, moral injury, error accumulation, or long-term fragility.

Proposition K.1 (Goodhart Collapse). *If an institution minimizes surprisal under a generative model whose likelihood depends only on $m \circ \pi_D$, then optimization pressure cannot distinguish beneficial improvements in X from destructive movements inside the same projected fiber. Consequently, sufficiently strong optimization generically converts projection blind spots into institutional damage.*

Proof. The likelihood depends only on $m(\pi_D(x))$. Therefore any two states x_1, x_2 with $m(\pi_D(x_1)) = m(\pi_D(x_2))$ are equivalent under the institutional model, regardless of how they differ in X . Optimization can move from x_1 to x_2 without penalty whenever the movement preserves or improves the metric. If the unmeasured coordinates of X include variables necessary for long-term admissibility, such movements can reduce true reachability while appearing successful under D . This is precisely the formal structure of Goodhart collapse. $\square$

The proposition sharpens the claim made informally in Section 5: a managing system does not need to intend harm for metric optimization to erode what it cannot see. Sufficient optimization pressure on $m \circ \pi_D$ is enough on its own.

L A Seam as a Hyperprior on Framework Change

This section restates the seam-expansion result of Appendix G probabilistically, giving a second formal reading of the claim from Section 8 that a seam keeps alternative frameworks live rather than foreclosed.

Let $D \in \mathcal{D}$ be a distinction system, and let a population maintain a distribution $P_t(D)$ over possible frameworks. A completely closed ontology is the degenerate distribution

$$P_t(D) = \delta_{D_t},$$

which assigns all probability to the currently active framework.

A visible seam prevents this collapse by maintaining nonzero probability on at least one alternative framework:

$$\exists D' \neq D_t \text{ such that } P_t(D') > 0.$$

In Bayesian language, the seam functions as a hyperprior over possible descriptions. It does not tell the population which alternative framework is correct; it merely prevents the prior over frameworks from collapsing to a point mass.

The entropy of the framework distribution,

$$H(P_t) = - \sum_{D \in \mathcal{D}} P_t(D) \log P_t(D),$$

measures one aspect of representation-shifting capacity. Removing seams drives $H(P_t)$ downward by eliminating live alternatives; adding seams weakly increases the support of P_t and can therefore increase the entropy of available descriptions.

Proposition L.1 (Seams Preserve Framework Entropy). *A visible seam preserves meta-ontological capacity by preventing the distribution over distinction systems from collapsing to a delta distribution. In the limiting case where all seams are removed, $P_t(D) \rightarrow \delta_{D_t}$ and framework entropy tends to zero.*

Proof. A seam makes at least one transition $D_t \rightarrow D'$ representable and available. Therefore the support of P_t contains more than one element whenever the seam is live and assigned nonzero probability. A delta distribution has zero entropy, while any distribution with support on more than one framework has nonnegative entropy and, when probabilities are nonzero on distinct alternatives, strictly positive entropy. Removing all seams removes all represented alternatives, forcing the support to collapse to $\{D_t\}$ and hence $H(P_t) \rightarrow 0$. $\square$

This gives a compact probabilistic restatement of the essay's main thesis. A domination ecology is not merely one that prefers bad answers. It is one that drives the prior over possible frameworks toward a delta distribution. A regenerative ecology keeps the prior over frameworks open.

Appendices J through L extend the core chain in three directions rather than adding to it linearly. The free energy principle is recovered as fixed-framework admissibility, showing that

an existing research tradition in theoretical neuroscience already describes the degenerate, single-framework case of the apparatus developed here. Goodhart collapse is recovered as free energy minimization under a thin projection, showing that the institutional pathology of Section 5 and the cognitive pathology of fixed-framework self-maintenance are instances of the same mechanism viewed at different scales. And the seam, treated everywhere else in this essay as a discrete operator, is recovered as a hyperprior that keeps the probability distribution over frameworks from collapsing to a point mass. None of these three results is required for the essay's central claims; each is offered as evidence that the same formal move – recursing reachability and admissibility from the space of states to the space of descriptions – continues to generate testable structure well past the point where it was first introduced.

References

- [1] Flyxion (2026). *The Ecology of Distinctions*. Working monograph.
- [2] Flyxion (2026). *Beyond the Cognitive Light Cone: Reachability, Admissibility, and the Geometry of Future Possibility*. Working paper.
- [3] Alexander, C., Ishikawa, S., & Silverstein, M. (1977). *A Pattern Language: Towns, Buildings, Construction*. Oxford University Press, New York.
- [4] Bourdieu, P. (1977). *Outline of a Theory of Practice*. Cambridge University Press, Cambridge.
- [5] Bruineberg, J., Dołęga, K., Dewhurst, J., & Baltieri, M. (2022). The emperor's new Markov blankets. *Behavioral and Brain Sciences*, 45, e183.
- [6] Doctorow, C. (2026). *The Reverse Centaur's Guide to Life After AI: How to Think About Artificial Intelligence – Before It's Too Late*. MCD / Farrar, Straus and Giroux, New York.
- [7] Friston, K. (2010). The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127–138.
- [8] Goodhart, C. A. E. (1975). Problems of monetary management: the U.K. experience. In *Papers in Monetary Economics*, Vol. I. Reserve Bank of Australia, Sydney.
- [9] Illich, I. (1973). *Tools for Conviviality*. Harper & Row, New York.
- [10] Kuhn, T. S. (1962). *The Structure of Scientific Revolutions*. University of Chicago Press, Chicago.
- [11] Luhmann, N. (1995). *Social Systems*. Stanford University Press, Stanford. (Translation of *Soziale Systeme*, 1984.)
- [12] Pearl, J. (1988). *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan Kaufmann, San Francisco, CA.

- [13] Strathern, M. (1997). 'Improving ratings': audit in the British University system. *European Review*, 5(3), 305–321.
- [14] Suber, P. (1990). *The Paradox of Self-Amendment: A Study of Law, Logic, Omnipotence, and Change*. Peter Lang, New York.