



Projection, Memory, and the Architecture of Faithful Representation
Flyxion, 2026.




The Core Tension: Capability vs. Faithfulness

SYSTEM CAPABILITY



High Prediction: The system matches the target output perfectly. The exterior is pristine; it achieves low perplexity and mastery over observational equivalence.

SYSTEM FAITHFULNESS



Low Faithfulness: The internal causal structure is broken. It predicts flawlessly for the wrong reasons. The representation has discarded the structure that actually matters for intervention and explanation.

What is the difference between a system that predicts correctly and a system that understands?

The Foundational Law: Every Representation is a Projection

The Curse of Complete Description

A representation that preserves every detail of reality possesses the complexity of reality itself. It ceases to be a model and becomes a duplicate.


The Projection Theorem

To reduce complexity, a representation must map multiple distinct states to the same output. It must collapse distinctions.

Structured Forgetting

Compression is not a defect of explanation; it is its enabling condition. The challenge is ensuring the forgotten information does not matter.

COMPLEXITY SCALE: 1:1




The Metric of Relevance

Admissible Distinctions

A distinction between two states only matters if it alters the outcome of an admissible task. A faithful representation preserves only the distinctions encoded by the admissibility manifold.

Ontological Deficit


A structural blind spot. Arises when the representational vocabulary is incapable of expressing a distinction at all (e.g., the Weierstrass obstruction, where smooth templates cannot capture recursive, scale-coupling structure). Invisible from within.



Epistemology is the theory of which *collapses* matter.


The Illusion of Accuracy: Observational vs. Interventional Equivalence

Observational Equivalence



Two systems appear identical today.
They generate the exact same admissible
observations.


Interventional Equivalence



Two systems react the exact same way to a change
tomorrow. Causation is the propagation of distinctions
through admissible transformations.

Diagnostic Conclusion: Observational faithfulness does not entail interventional faithfulness.
A model can achieve observational equivalence while suffering massive causal distortion.


The Epistemic Crisis: Fixed Codecs vs. Living Archives



“A codec is a promise to forget nothing while understanding nothing. An archive that rematches promises to understand more, so it can store less.”





The Pressure of Ignorance: Residual Structure Density (Φ)



A template library under pressure from a concentrated Φ field is in an unstable equilibrium.

The Dynamics of Progress: Waves of Collapse




Mechanism: Minimum Description Length

New templates are admitted speculatively.


They earn their place—and trigger a collapse—only if they explain more residual mass than their own geometric cost.

Compression occurs in waves because understanding occurs in waves.

Anatomy of a Revolution: Distinction Reorganization



The Phase Shift: The World-Model Crossover




Core Principle: The boundary between data collection and theoretical science is mathematically definable.

The Unit of Surprise & Modality Integration

Diversity Requirement

The crossover threshold is not reached by accumulating more of the same data. It requires diversity of kernels—orthogonal observations (visual, acoustic, geometric) triangulating the same domain.



Cross-Modal Cascades


An audio recording of an echo can retroactively compress a photograph taken eight years earlier by refining the shared 3D geometry template. Sound borrows geometry.

The Byte: Storage cost no longer tracks the data rate of the camera; it asymptotically tracks the novelty rate of the world. **The byte is the unit of surprise.**



Application I: Memory as Generative Reconstruction



1. 2018: Ingestion
(Raw sensor data leaves a partial residual record).



2. 2026: The Evolving Library
(The distinction structure matures, creating new templates).



3. Rendered Recall
(The past is rendered through the vocabulary of the present).


Key Insight

Memory is not playback; it is rendering. We do not retrieve an old file. We **render** an old residual through our most recent understanding of the world.

Representational Revision Principle


Distinction reorganization modifies the representation of prior observations, not the observations themselves. The past becomes more compressible over time.

Application II: Interpretability as Distinction Verification



The Description Fallacy

A model prompted to produce chain-of-thought reasoning may fluently describe a trajectory without actually performing it. Coherent description can be post-hoc confabulation.




The Witness Principle

A claim that a distinction is represented in a model requires a constructive Recovery Operator. Interpretability is not describing internal structure; it is verifying whether the model preserves the admissible distinctions required for intervention.


Takeaway: Without a recovery operator, an interpretability claim is an assertion, not a proof.

Synthesis: Six Programs, One Geometry



Compression is not what we do to knowledge.
Compression is what knowledge looks like from the storage layer.

**Understanding is not the
accumulation of descriptions,
but a trajectory of distinction
reorganization: $D_0 \rightarrow D_1 \rightarrow D_2 \dots$**



**The geometry of understanding is not the geometry of facts in a space,
but the geometry of distinctions in transformation.**